

Loan_Payback

September 6, 2020

0.1 Kaushal Rao - Predicting Loan Payback with Decision Tree & Random Forest

0.1.1 In this notebook, we will try to predict which people have a high probability of paying an investor back with the Decision Tree and Random Forest statistical models. We will be exploring publicly available lending data from 2007-2010 (can be found on LendingClub.com).

Here are what the columns of the dataset represent: * credit.policy: 1 if the customer meets the credit underwriting criteria of LendingClub.com, and 0 otherwise. * purpose: The purpose of the loan (takes values "credit_card", "debt_consolidation", "educational", "major_purchase", "small_business", and "all_other"). * int.rate: The interest rate of the loan, as a proportion (a rate of 11% would be stored as 0.11). Borrowers judged by LendingClub.com to be more risky are assigned higher interest rates. * installment: The monthly installments owed by the borrower if the loan is funded. * log.annual.inc: The natural log of the self-reported annual income of the borrower. * dti: The debt-to-income ratio of the borrower (amount of debt divided by annual income). * fico: The FICO credit score of the borrower. * days.with.cr.line: The number of days the borrower has had a credit line. * revol.bal: The borrower's revolving balance (amount unpaid at the end of the credit card billing cycle). * revol.util: The borrower's revolving line utilization rate (the amount of the credit line used relative to total credit available). * inq.last.6mths: The borrower's number of inquiries by creditors in the last 6 months. * delinq.2yrs: The number of times the borrower had been 30+ days past due on a payment in the past 2 years. * pub.rec: The borrower's number of derogatory public records (bankruptcy filings, tax liens, or judgments).

0.1.2 Importing

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix
```

```
In [2]: loans = pd.read_csv('loan_data.csv')
```

```
In [3]: loans.describe()
```

```

Out [3]:      credit.policy      int.rate  installment  log.annual.inc      dti  \
count      9578.000000  9578.000000  9578.000000      9578.000000  9578.000000
mean         0.804970      0.122640    319.089413      10.932117    12.606679
std          0.396245      0.026847    207.071301       0.614813     6.883970
min          0.000000      0.060000     15.670000       7.547502     0.000000
25%          1.000000      0.103900    163.770000      10.558414     7.212500
50%          1.000000      0.122100    268.950000      10.928884    12.665000
75%          1.000000      0.140700    432.762500      11.291293    17.950000
max          1.000000      0.216400    940.140000      14.528354    29.960000

           fico  days.with.cr.line      revol.bal  revol.util  \
count      9578.000000      9578.000000  9.578000e+03  9578.000000
mean       710.846314      4560.767197  1.691396e+04   46.799236
std        37.970537      2496.930377  3.375619e+04   29.014417
min        612.000000      178.958333  0.000000e+00    0.000000
25%        682.000000      2820.000000  3.187000e+03   22.600000
50%        707.000000      4139.958333  8.596000e+03   46.300000
75%        737.000000      5730.000000  1.824950e+04   70.900000
max        827.000000      17639.958330  1.207359e+06  119.000000

           inq.last.6mths  delinq.2yrs      pub.rec  not.fully.paid
count      9578.000000  9578.000000  9578.000000      9578.000000
mean         1.577469      0.163708      0.062122      0.160054
std          2.200245      0.546215      0.262126      0.366676
min          0.000000      0.000000      0.000000      0.000000
25%          0.000000      0.000000      0.000000      0.000000
50%          1.000000      0.000000      0.000000      0.000000
75%          2.000000      0.000000      0.000000      0.000000
max          33.000000     13.000000      5.000000      1.000000

```

```

In [4]: loans.head()

```

```

Out [4]:      credit.policy      purpose  int.rate  installment  log.annual.inc  \
0           1  debt_consolidation    0.1189      829.10      11.350407
1           1    credit_card        0.1071      228.22      11.082143
2           1  debt_consolidation    0.1357      366.86      10.373491
3           1  debt_consolidation    0.1008      162.34      11.350407
4           1    credit_card        0.1426      102.92      11.299732

           dti  fico  days.with.cr.line  revol.bal  revol.util  inq.last.6mths  \
0    19.48    737      5639.958333      28854      52.1          0
1    14.29    707      2760.000000      33623      76.7          0
2    11.63    682      4710.000000       3511      25.6          1
3     8.10    712      2699.958333      33667      73.2          1
4    14.97    667      4066.000000       4740      39.5          0

           delinq.2yrs  pub.rec  not.fully.paid
0           0          0          0

```


1	0	0	0
2	0	0	0
3	0	0	0
4	1	0	0

0.1.3 Exploratory Data Analysis (EDA)

```
In [5]: plt.figure(figsize=(10,6))
        loans[loans['credit.policy']==1]['fico'].hist(alpha=0.5,color='blue',
                                                    bins=30,label='Credit.Policy=1')
        loans[loans['credit.policy']==0]['fico'].hist(alpha=0.5,color='red',
                                                    bins=30,label='Credit.Policy=0')


        plt.legend()
        plt.xlabel('FICO')
        # creating histogram of the two FICO distributions, one for each credit policy outcome
        # we see that those that have met the credit requirements generally have a higher FICO
```

Out [5]: Text(0.5, 0, 'FICO')




```
In [6]: plt.figure(figsize=(11,7))
        sns.countplot(x='purpose',hue='not.fully.paid',data=loans,palette='Set1')
        # countplot showing the counts of loans by purpose
```

Out [6]: <matplotlib.axes._subplots.AxesSubplot at 0x103f012e8>



```
In [7]: sns.jointplot(x='fico',y='int.rate',data=loans,color='purple')  
        # jointplot between FICO scores and interest rate  
        # generally inverse correlation
```


```
Out[7]: <seaborn.axisgrid.JointGrid at 0x1a182c45c0>
```



```
In [8]: plt.figure(figsize=(11,7))
sns.lmplot(y='int.rate',x='fico',data=loans,hue='credit.policy',
           col='not.fully.paid',palette='Set1')
# linear model plots to see trends in FICO scores based on "fully paid" and credit pol
```

```
Out[8]: <seaborn.axisgrid.FacetGrid at 0x1a1858aba8>
```

```
<Figure size 792x504 with 0 Axes>
```



0.1.4 Setting up the Data

In [9]: `loans.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
credit.policy      9578 non-null int64
purpose           9578 non-null object
int.rate          9578 non-null float64
installment       9578 non-null float64
log.annual.inc    9578 non-null float64
dti               9578 non-null float64
fico              9578 non-null int64
days.with.cr.line 9578 non-null float64
revol.bal         9578 non-null int64
revol.util        9578 non-null float64
inq.last.6mths    9578 non-null int64
delinq.2yrs       9578 non-null int64
pub.rec           9578 non-null int64
not.fully.paid    9578 non-null int64
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```

Categorical Features

```
In [10]: # since the "purpose" column is categorical, we need to transform it
         # we can use dummy variables so sklearn is able to understand it as a feature
         cat_feats = ['purpose']
```

```

final_data = pd.get_dummies(loans,columns=cat_feats,drop_first=True)
final_data.info()
# we can see that the purpose binary columns have been added

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 19 columns):
credit.policy          9578 non-null int64
int.rate               9578 non-null float64
installment           9578 non-null float64
log.annual.inc        9578 non-null float64
dti                   9578 non-null float64
fico                  9578 non-null int64
days.with.cr.line    9578 non-null float64
revol.bal             9578 non-null int64
revol.util            9578 non-null float64
inq.last.6mths        9578 non-null int64
delinq.2yrs           9578 non-null int64
pub.rec              9578 non-null int64
not.fully.paid        9578 non-null int64
purpose_credit_card   9578 non-null uint8
purpose_debt_consolidation 9578 non-null uint8
purpose_educational   9578 non-null uint8
purpose_home_improvement 9578 non-null uint8
purpose_major_purchase 9578 non-null uint8
purpose_small_business 9578 non-null uint8
dtypes: float64(6), int64(7), uint8(6)
memory usage: 1.0 MB

```

Train-Test Split

```

In [11]: X = final_data.drop('not.fully.paid',axis=1)
         y = final_data['not.fully.paid']
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=42)

```

0.1.5 Decision Tree

Training

```

In [12]: dtree = DecisionTreeClassifier()
         dtree.fit(X_train,y_train)

Out[12]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                splitter='best')

```

Predictions & Evaluation

```
In [13]: predictions = dtree.predict(X_test)
```

```
In [14]: print(classification_report(y_test,predictions))
         # not great for predicting class 1!
```

	precision	recall	f1-score	support
0	0.86	0.82	0.84	2431
1	0.20	0.24	0.22	443
micro avg	0.73	0.73	0.73	2874
macro avg	0.53	0.53	0.53	2874
weighted avg	0.75	0.73	0.74	2874

```
In [15]: print(confusion_matrix(y_test,predictions))
```

```
[[2001  430]
 [ 336  107]]
```

0.1.6 Random Forest

Training

```
In [16]: rfc = RandomForestClassifier(n_estimators=600)
         rfc.fit(X_train,y_train)
```

```
Out[16]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=600, n_jobs=None,
                                oob_score=False, random_state=None, verbose=0,
                                warm_start=False)
```

Predictions & Evaluation

```
In [17]: predictions = rfc.predict(X_test)
```

```
In [18]: print(classification_report(y_test,predictions))
         # better than decision tree but still not great for class 1
```

	precision	recall	f1-score	support
0	0.85	1.00	0.92	2431
1	0.53	0.02	0.04	443

micro avg	0.85	0.85	0.85	2874
macro avg	0.69	0.51	0.48	2874
weighted avg	0.80	0.85	0.78	2874

```
In [19]: print(confusion_matrix(y_test,predictions))
          # recall was not great for predicting class 1, more feature engineering is needed
```

```
[[2422   9]
 [ 433 10]]
```